

Mohan Zhang

mohanz@cs.unc.edu • [Homepage](#)

RESEARCH INTEREST

Efficient and Reliable ML

- Building intelligent agent systems for web navigation, robotic automation, and other real-world applications.
- Advancing the scalability, efficiency, and reliability of LLMs, with interests spanning MoE, sparsification, model auditing, and adversarial robustness.

EDUCATION

University of North Carolina at Chapel Hill, Chapel Hill, NC, USA

Aug 2025 – Present

- Ph.D. in Computer Science
 - Advisor: [Prof. Tianlong Chen](#)

University of Science and Technology of China, Hefei, Anhui, China

Sep 2021 – Jun 2025

- B.Eng. in Computer Science and Technology (GPA: 3.85/4.3)

PUBLICATIONS (*: Equal contribution)

[TRUST: A Crowdsourcing Framework for Auditing Large Language Model Reasoning](#)

Morris Yu-Chao Huang, Zhen Tan, **Mohan Zhang**, Pingzhi Li, Zhuo Zhang, Tianlong Chen
COLM 2026

[Measuring Real-World Prompt Injection Attacks in LLM-based Resume Screening](#)

Mohan Zhang, Yuqi Jia, Zhen Tan, Steven Jiang, Neil Zhenqiang Gong, Tianlong Chen, Dawn Song
35th USENIX Security Symposium, 2026

[Bridging Language and Physics: Automated Design of Continuum Robots with Large Language Models](#)

Jingyi Chen*, **Mohan Zhang***, Laura Yao*, Yingtai Ni, Jianmin Ji, Jie Peng, Song Wang, Tianlong Chen
RSS 2026

[Symbiotic Cooperation for Web Agents: Harnessing Complementary Strengths of Large and Small LLMs](#)

Ruichen Zhang*, Mufan Qiu*, Zhen Tan*, **Mohan Zhang**, Xiaopeng Lu, Jie Peng, Kaidi Xu, Leandro Z. Agudelo, Peter Zhenghao Qian, Tianlong Chen
CPAL 2026

[One Token Embedding Is Enough to Deadlock Your Large Reasoning Model](#)

Mohan Zhang*, Yihua Zhang*, Jinghan Jia, Zhangyang “Atlas” Wang, Sijia Liu, Tianlong Chen
NeurIPS 2025

[Modalities Contribute Unequally: Enhancing Medical Multi-modal Learning through Adaptive Modality Token Re-balancing](#)

Jie Peng*, Jenna L. Ballard*, **Mohan Zhang**, Sukwon Yun, Jiayi Xin, Qi Long, Yanyong Zhang, Tianlong Chen
ICML 2025

[Advancing MoE Efficiency: A Collaboration-Constrained Routing \(C2R\) Strategy for Better Expert Parallelism Design](#)

Mohan Zhang*, Pingzhi Li*, Jie Peng, Mufan Qiu, Tianlong Chen
NAACL 2025; SAC Award for Low-resource Methods for NLP

IN SUBMISSION

[Tx-Bench: A Benchmark for Open-Ended Ethereum Transaction Annotation](#)

Mohan Zhang, Yu-Chao Huang, Jie Peng, Zhuo Zhang, Tianlong Chen
Submitted to EMNLP 2026

[ASSET: Acquisition-Sensitive Subspace Estimation for Test-Time Adaptation of Medical VLMs](#)

Mohan Zhang, Zhen Tan, Songyuan Sui, Tianlong Chen
Submitted to NeurIPS 2026

[Causal Audit and Active Refinement for Multi-Agent Swarms via Graph Reconstruction](#)

Yu-Chao Huang, Pingzhi Li, Jie Peng, **Mohan Zhang**, Zhuo Zhang, Tianlong Chen
Submitted to NeurIPS 2026

[Do More Modalities Always Help? A Geometric Perspective on Missing-Modality Robustness](#)

Songyuan Sui, Zhen Tan, **Mohan Zhang**, Rana Muhammad Shahroz Khan, Xia Hu, Tianlong Chen
Submitted to NeurIPS 2026

PPBoost: Progressive Prompt Boosting for Text-Driven Medical Image Segmentation

Xuchen Li, Hengrui Gu, **Mohan Zhang**, Qin Liu, Zhen Tan, Xinyuan Zhu, Huixue Zhou, Tianlong Chen, Kaixiong Zhou
Submitted to NeurIPS 2026

RESEARCH EXPERIENCE	University of North Carolina at Chapel Hill, Chapel Hill, NC	Feb 2024 – Jun 2025
	▪ Remote Research Intern, mentored by Prof. Tianlong Chen	
HONORS & AWARDS	▪ Low-resource Methods for NLP SAC Award, NAACL 2025	2025
	▪ Outstanding Student Scholarship, University of Science and Technology of China	2022, 2023, 2024